Responsible AI Statement

Verifai is designed strictly as an assistive tool to evaluate digital media provenance and surface structural probability fields. Because machine learning systems possess inherent statistical error margins, false classifications present clear operational risks, including mislabeling authentic content as synthetic or validating deepfakes. Consequently, this engine is explicitly out-of-scope for high-stakes decision-making, including legal evidence verification, automated content moderation pipelines, or automated judicial determinations where absolute, unyielding deterministic certainty is required.

To manage model uncertainty responsibly, Verifai rejects forced binary labels. By exposing raw confidence scores across strict Pydantic thresholds, we explicitly flag classification ambiguity. Users are instructed to treat any probability output landing within the [0.26 - 0.75] threshold range as an inconclusive evaluation requiring manual human verification.

Finally, traditional media forensics force users to sacrifice data sovereignty by uploading private imagery to centralized cloud servers. Verifai natively mitigates this risk through a privacy-first, local architecture. Because 100% of tensor extraction, Fast Fourier Transformation, and ONNX model execution occurs directly within the client's local browser memory, user imagery data is never transmitted, cached, or compiled by third parties, eliminating server-side surveillance risks entirely.